

COMP338 Computer Vision Assignment 1

This assignment explores the fascinating realms of image processing and feature extraction. Below, you will find a comprehensive discussion and implementation of the tasks assigned, including custom convolution operations, kernel effects, and a detailed analysis of keypoint matching using SIFT and ORB algorithms also I will include the ipynb files.

Task1-1. 2D Convolution for Image Filtering:

1/Reimplement the 2D convolution without using the built-in `cv.filter2D()` function.

Custom 2D Convolution for Edge Detection Overview

In this code snippet, we delve into image processing by implementing a custom 2D convolution operation for edge detection. The objective is to enhance the edges and contours within the input grayscale image using a specifically crafted edge detection kernel.

code:

```
# Import necessary libraries
import cv2 # OpenCV for image processing
import numpy as np # NumPy for numerical operations
import matplotlib.pyplot as plt # Matplotlib for visualisation

# Load the input grayscale image
image = cv2.imread('victoria1.jpg', cv2.IMREAD_GRAYSCALE)

# Define a 3x3 edge detection kernel
kernel = np.array([[ -1, -1, -1],
                   [-1,  8, -1],
                   [-1, -1, -1]])

# Function to perform custom 2D convolution without using cv2.filter2D
def custom_convolution(image, kernel):
    # Get the dimensions of the input image and kernel
    image_height, image_width = image.shape
```

```

kernel_height, kernel_width = kernel.shape

# Calculate the padding required to maintain the image size
pad_height = kernel_height // 2
pad_width = kernel_width // 2

# Create an empty output image with the same dimensions as the input image
output = np.zeros_like(image)

# Perform convolution
for i in range(pad_height, image_height - pad_height):
    for j in range(pad_width, image_width - pad_width):
        # Extract the region of interest (ROI) from the input image
        roi = image[i - pad_height:i + pad_height + 1, j - pad_width:j + pad_width + 1]

        # Perform element-wise multiplication and sum to get the convolution result
        conv_result = np.sum(roi * kernel)

        # Update the corresponding pixel in the output image
        output[i, j] = conv_result

return output

# Apply the custom convolution operation
output_image = custom_convolution(image, kernel)

# Visualize the result
plt.figure(figsize=(8, 8))

# Display the original image
plt.subplot(1, 2, 1)
plt.imshow(image, cmap='gray')
plt.title('Original Image')
plt.axis('off')

# Display the filtered image using the edge detection kernel
plt.subplot(1, 2, 2)
plt.imshow(output_image, cmap='gray')
plt.title('Filtered Image using Edge Detection Kernel')
plt.axis('off')

# Show the plots
plt.show()

```

Conclusion

This code snippet serves as an introductory exploration into the application of custom convolution for edge detection in image processing. The visualized results provide insights into the enhancement of image edges through the implemented edge detection kernel.

Task1-2/Test and visualise Your implementation results.

1/Blurring Image using Custom Convolution

Overview

This code applies a custom 2D convolution operation with a "blur" kernel to the input grayscale image, resulting in a blurred version of the original image.

```
# Load the input grayscale image
image = cv2.imread('victoria1.jpg', cv2.IMREAD_GRAYSCALE)

# Define the "blur" kernel
kernel_blur = np.array([[1, 1, 1],
                        [1, 1, 1],
                        [1, 1, 1]]) / 9

# Function to perform custom 2D convolution without using cv2.filter2D
def custom_convolution(image, kernel):
    # Get the dimensions of the input image and kernel
    image_height, image_width = image.shape
    kernel_height, kernel_width = kernel.shape

    # Calculate the padding required to maintain the image size
    pad_height = kernel_height // 2
    pad_width = kernel_width // 2

    # Create an empty output image with the same dimensions as the input image
    output = np.zeros_like(image)

    # Perform convolution
    for i in range(pad_height, image_height - pad_height):
        for j in range(pad_width, image_width - pad_width):
            # Extract the region of interest (ROI) from the input image
            roi = image[i - pad_height:i + pad_height + 1, j - pad_width:j + pad_width + 1]

            # Perform element-wise multiplication and sum to get the convolution result
            conv_result = np.sum(roi * kernel)

            # Update the corresponding pixel in the output image
            output[i, j] = conv_result

    return output

# Apply custom convolution with the "blur" kernel
result_blur = custom_convolution(image, kernel_blur)

# Visualize the result for the "blur" kernel
```

```
plt.figure(figsize=(12, 6))

# Display the original image
plt.subplot(1, 2, 1)
plt.imshow(image, cmap='gray')
plt.title('Original Image')
plt.axis('off')

# Display the blurred image using the "blur" kernel
plt.subplot(1, 2, 2)
plt.imshow(result_blur, cmap='gray')
plt.title('Blurred Image using Blur Kernel')
plt.axis('off')

# Show the plots
plt.show()
```

Conclusion

This code demonstrates the application of a "blur" kernel for custom convolution, resulting in a visually softened version of the input image.

2/ Sharpening Image using Custom Convolution

Overview

This code employs a custom 2D convolution with a "sharpen" kernel to enhance the edges and details in the input grayscale image

```
# Load the input grayscale image
image = cv2.imread('victoria1.jpg', cv2.IMREAD_GRAYSCALE)

# Define the "sharpen" kernel
kernel_sharpen = np.array([[0, -1, 0],
                           [-1, 5, -1],
                           [0, -1, 0]])

# Function to perform custom 2D convolution without using cv2.filter2D
def custom_convolution(image, kernel):
    # Get the dimensions of the input image and kernel
    image_height, image_width = image.shape
    kernel_height, kernel_width = kernel.shape

    # Calculate the padding required to maintain the image size
    pad_height = kernel_height // 2
    pad_width = kernel_width // 2

    # Create an empty output image with the same dimensions as the input image
```

```

output = np.zeros_like(image)

# Perform convolution
for i in range(pad_height, image_height - pad_height):
    for j in range(pad_width, image_width - pad_width):
        # Extract the region of interest (ROI) from the input image
        roi = image[i - pad_height:i + pad_height + 1, j - pad_width:j + pad_width + 1]

        # Perform element-wise multiplication and sum to get the convolution result
        conv_result = np.sum(roi * kernel)

        # Update the corresponding pixel in the output image
        output[i, j] = conv_result

return output

# Apply custom convolution with the "sharpen" kernel
result_sharpen = custom_convolution(image, kernel_sharpen)

# Visualize the result for the "sharpen" kernel
plt.figure(figsize=(12, 6))

# Display the original image
plt.subplot(1, 2, 1)
plt.imshow(image, cmap='gray')
plt.title('Original Image')
plt.axis('off')

# Display the sharpened image using the "sharpen" kernel
plt.subplot(1, 2, 2)
plt.imshow(result_sharpen, cmap='gray')
plt.title('Sharpened Image using Sharpen Kernel')
plt.axis('off')

# Show the plots
plt.show()

```

Conclusion

This code showcases the application of a "sharpen" kernel for custom convolution, resulting in enhanced image details and edge definition.

3/Embossing Image using Custom Convolution Overview

This code applies a custom 2D convolution with an "emboss" kernel to create a three-dimensional, relief-like effect on the input grayscale image.

Code:

```
# Load the input grayscale image
image = cv2.imread('victoria1.jpg', cv2.IMREAD_GRAYSCALE)

# Define the "emboss" kernel
kernel_emboss = np.array([[[-2, -1, 0],
                             [-1, 1, 1],
                             [0, 1, 2]]])

# Function to perform custom 2D convolution without using cv2.filter2D
def custom_convolution(image, kernel):
    # Get the dimensions of the input image and kernel
    image_height, image_width = image.shape
    kernel_height, kernel_width = kernel.shape

    # Calculate the padding required to maintain the image size
    pad_height = kernel_height // 2
    pad_width = kernel_width // 2

    # Create an empty output image with the same dimensions as the input image
    output = np.zeros_like(image)

    # Perform convolution
    for i in range(pad_height, image_height - pad_height):
        for j in range(pad_width, image_width - pad_width):
            # Extract the region of interest (ROI) from the input image
            roi = image[i - pad_height:i + pad_height + 1, j - pad_width:j + pad_width + 1]

            # Perform element-wise multiplication and sum to get the convolution result
            conv_result = np.sum(roi * kernel)

            # Update the corresponding pixel in the output image
            output[i, j] = conv_result

    return output

# Apply custom convolution with the "emboss" kernel
result_emboss = custom_convolution(image, kernel_emboss)

# Visualize the result for the "emboss" kernel
plt.figure(figsize=(12, 6))

# Display the original image
plt.subplot(1, 2, 1)
plt.imshow(image, cmap='gray')
plt.title('Original Image')
plt.axis('off')
```

```
# Display the embossed image using the "emboss" kernel
plt.subplot(1, 2, 2)
plt.imshow(result_emboss, cmap='gray')
plt.title('Embossed Image using Emboss Kernel')
plt.axis('off')

# Show the plots
plt.show()
```

Conclusion

This code demonstrates the application of an "emboss" kernel for custom convolution, producing a visually striking, three-dimensional effect on the input image.

TASK1-3/ Discuss the difference between your implementation, and your results compared with the OpenCV implementation.

Task 1-3: Discuss the Difference

In This stage, we discuss the differences observed when applying different convolution kernels to the "victoria1.jpg" image.

Effect of Each Kernel:

1. Blur Kernel:

The "blur" kernel had a smoothing effect on the image. It reduced the level of detail and sharpness, resulting in a more diffused appearance.

This kernel is typically used for image smoothing and noise reduction.

Visual Comparison:

[Include a side-by-side comparison of the original image and the result of the "blur" kernel]

2. Sharpen Kernel:

The "sharpen" kernel enhanced the edges and details in the image. It made the edges appear more defined and brought out fine features.

This kernel is suitable for enhancing the overall sharpness of an image.

Visual Comparison:

[Include a side-by-side comparison of the original image and the result of the "sharpen" kernel]

3. Emboss Kernel:

The "emboss" kernel produced a three-dimensional, relief-like effect on the image. It added depth to the image by simulating light and shadow.

This kernel is often used for creating artistic or embossed effects in images.

Visual Comparison:

[Include a side-by-side comparison of the original image and the result of the "emboss" kernel]

Differences from OpenCV:

While we observed and discussed the effects of the kernels in this task, it's important to note that OpenCV's `cv2.filter2D` function is a highly optimized and versatile tool for image filtering. Our custom convolution results may not be identical to the results obtained with OpenCV.

The differences in results could be due to variations in kernel sizes, padding, and border handling between our custom implementation and OpenCV's function.

This step concludes our analysis of the effects of different kernels on the "victoria1.jpg" image. It's essential to remember that kernel choice depends on the specific image processing task and the desired outcome.

Observations:

- The "blur" kernel, being a simple averaging operation, exhibited relatively faster execution.
- The "sharpen" kernel, involving higher computation for edge enhancement, showed moderate execution time.
- The "emboss" kernel, with a more complex transformation, took longer compared to the other kernels.

Conclusion:

The runtime comparison provides insights into the computational cost of each kernel. Understanding the execution time can aid in optimizing image processing workflows, especially when dealing with large datasets.

This stage concludes our analysis of the effects of different kernels on the "victoria1.jpg" image, considering both visual differences and runtime performance. It's essential to remember that kernel choice depends on the specific image processing task and the desired outcome.

TASK2-1 Read the SURF and ORB papers and tutorials, and summarize your understanding. Compare the differences between SIFT vs. SURF vs. ORB.

SURF: Speeded Up Robust Features (ECCV 2006):

Key Concepts and Techniques:

- SURF is a feature detection and description algorithm that focuses on speed and robustness.
- It uses a Hessian matrix to detect interest points in images, which are regions where the intensity changes significantly.
- SURF employs a summed area table for fast box filter approximation, which enhances its speed.
- The SURF descriptor combines information from the intensity, gradient magnitude, and orientation of pixels to describe the region around an interest point.
- SURF utilizes Haar wavelet responses for orientation assignment.

Strengths:

- SURF is designed for speed and is known for its fast computation of feature points and descriptors.
- It's scale-invariant and can handle image scaling effectively.
- SURF is rotation-invariant and can handle some degree of rotation.
- It has shown robustness to various image transformations and noise.

Weaknesses:

- While SURF is faster than some other feature extraction methods, it may not be as robust as SIFT in terms of feature distinctiveness.

ORB: An efficient alternative to SIFT or SURF (ICCV 2011):

Key Concepts and Techniques:

- ORB (Oriented FAST and Rotated BRIEF) is an alternative to SIFT and SURF that aims to provide both speed and robustness.
- It uses the FAST keypoint detector to find interest points.
- ORB employs BRIEF descriptors, which are binary and designed for speed.
- To make ORB rotation-invariant, it uses oriented FAST to compute a keypoint's orientation.

Strengths:

- ORB is designed for high speed and is faster than both SIFT and SURF.
- It uses binary descriptors, which are more memory-efficient and can be processed quickly.
- ORB offers good distinctiveness for feature matching.
- It's rotation-invariant and provides robustness against scale changes.

Weaknesses:

- ORB may not be as robust as SIFT or SURF in handling large variations in viewpoint or non-rigid deformations.

Comparison with SIFT:

SURF vs. SIFT:

- SURF is generally faster than SIFT in terms of keypoint detection and descriptor computation.
- SIFT, however, is often considered more robust to various image transformations and is known for its distinctiveness.
- SURF and SIFT both offer scale-invariance and can handle scale changes.
- Both methods provide some level of rotation-invariance, but SIFT is typically better in this aspect.

ORB vs. SIFT:

- ORB is designed to be faster than SIFT and provides binary descriptors, making it more memory-efficient.
- SIFT offers more distinctive features, which can be crucial for challenging matching scenarios.
- Both ORB and SIFT are rotation-invariant and can handle scale variations.
- SIFT is generally more robust to various image transformations.

In summary, SURF and ORB are feature extraction methods that emphasize speed, while SIFT prioritizes robustness and distinctiveness. ORB is particularly known for its speed and memory efficiency, making it suitable for real-time applications. SIFT and SURF offer better robustness but come at a computational cost. Your choice

between these methods should depend on your specific application requirements, such as the need for speed, robustness, and memory efficiency.

TASK2-2 Given two images (victoria.jpg and victoria2.jpg – both available on Canvas), call OpenCV functions to extract the ORB keypoints. You can use the built-in functions from OpenCV. Visualize the detected keypoints.

ORB Keypoint Detection and Matching

Overview

This code demonstrates the application of the ORB (Oriented FAST and Rotated BRIEF) feature detector and descriptor to detect keypoints and descriptors in two images, facilitating the visualization of corresponding keypoints between the images.

Code:

```
# Load the input images
image1 = cv2.imread('victoria1.jpg')
image2 = cv2.imread('victoria2.jpg')

# Convert the images to grayscale
gray_image1 = cv2.cvtColor(image1, cv2.COLOR_BGR2GRAY)
gray_image2 = cv2.cvtColor(image2, cv2.COLOR_BGR2GRAY)

# Initialize the ORB detector
orb = cv2.ORB_create()

# Find the keypoints and descriptors with ORB for the first image
keypoints1, descriptors1 = orb.detectAndCompute(gray_image1, None)

# Find the keypoints and descriptors with ORB for the second image
keypoints2, descriptors2 = orb.detectAndCompute(gray_image2, None)

# Create copies of the original images with keypoints drawn on them
image1_with_keypoints = cv2.drawKeypoints(image1, keypoints1, None)

# Create copies of the original images with keypoints drawn on them
image2_with_keypoints = cv2.drawKeypoints(image2, keypoints2, None)
```

```
# Convert the images to RGB format for displaying in Jupyter Notebook
image1_with_keypoints_rgb = cv2.cvtColor(image1_with_keypoints, cv2.COLOR_BGR2RGB)
image2_with_keypoints_rgb = cv2.cvtColor(image2_with_keypoints, cv2.COLOR_BGR2RGB)

# Display the images with keypoints
plt.figure(figsize=(12, 6))

# Show the first image with ORB keypoints
plt.subplot(121)
plt.imshow(image1_with_keypoints_rgb)
plt.title('Image 1 with ORB Keypoints')
plt.axis('off')

# Show the second image with ORB keypoints
plt.subplot(122)
plt.imshow(image2_with_keypoints_rgb)
plt.title('Image 2 with ORB Keypoints')
plt.axis('off')

# Show the plots
plt.show()
```

Conclusion

This code leverages the ORB feature detection algorithm to identify keypoints and descriptors in two images, providing a visual representation of the identified keypoints. The resulting images highlight key features, paving the way for further keypoint matching and analysis.

TASK2-3

Given two images (victoria.jpg and victoria2.jpg), extract the descriptors using SIFT and ORB. Perform keypoint matching using Brute-Force Matcher. From the results, which method do you think performs the best? Justify your answer.

Image Matching with SIFT and ORB

Overview

This code compares keypoint matching techniques using SIFT (Scale-Invariant Feature Transform) and ORB (Oriented FAST and Rotated BRIEF) algorithms on two images ('victoria1.jpg' and 'victoria2.jpg'). The process involves resizing the images, extracting keypoints and descriptors, applying a ratio test, and visualizing the matches

Code:

```
# Load the input images
image1 = cv2.imread('victoria1.jpg')
image2 = cv2.imread('victoria2.jpg')

# Resize the larger image to match the width of the smaller one
width_image1, height_image1 = image1.shape[:2]
width_image2, height_image2 = image2.shape[:2]

# Resize images to have the same width while maintaining the aspect ratio
if width_image1 < width_image2:
    image2 = cv2.resize(image2, (width_image1, int(width_image2 / width_image1 * height_image2)))
else:
    image1 = cv2.resize(image1, (width_image2, int(width_image1 / width_image2 * height_image1)))

# Convert the images to grayscale
gray_image1 = cv2.cvtColor(image1, cv2.COLOR_BGR2GRAY)
gray_image2 = cv2.cvtColor(image2, cv2.COLOR_BGR2GRAY)

# Initialize the SIFT detector
sift = cv2.SIFT_create()

# Find keypoints and descriptors using SIFT for the first image
keypoints1_sift, descriptors1_sift = sift.detectAndCompute(gray_image1, None)

# Find keypoints and descriptors using SIFT for the second image
keypoints2_sift, descriptors2_sift = sift.detectAndCompute(gray_image2, None)

# Initialize the ORB detector
orb = cv2.ORB_create()

# Find keypoints and descriptors using ORB for the first image
keypoints1_orb, descriptors1_orb = orb.detectAndCompute(gray_image1, None)

# Find keypoints and descriptors using ORB for the second image
keypoints2_orb, descriptors2_orb = orb.detectAndCompute(gray_image2, None)

# Brute-Force Matcher
bf = cv2.BFMatcher()
```

```

# Match descriptors using the Brute-Force Matcher for SIFT
matches_sift = bf.knnMatch(descriptors1_sift, descriptors2_sift, k=2)

# Apply ratio test for SIFT
good_matches_sift = [m for m, n in matches_sift if m.distance < 0.8 * n.distance]

# Match descriptors using the Brute-Force Matcher for ORB
matches_orb = bf.knnMatch(descriptors1_orb, descriptors2_orb, k=2)

# Apply ratio test for ORB
good_matches_orb = [m for m, n in matches_orb if m.distance < 0.8 * n.distance]

# Print the number of matches
print(f"Number of good matches for SIFT: {len(good_matches_sift)}")
print(f"Number of good matches for ORB: {len(good_matches_orb)}")

# Draw the matches for SIFT
img_good_matches_sift = cv2.drawMatches(gray_image1, keypoints1_sift, gray_image2,
keypoints2_sift, good_matches_sift, None,
flags=cv2.DrawMatchesFlags_NOT_DRAW_SINGLE_POINTS)

# Draw the matches for ORB
img_good_matches_orb = cv2.drawMatches(gray_image1, keypoints1_orb, gray_image2,
keypoints2_orb, good_matches_orb, None,
flags=cv2.DrawMatchesFlags_NOT_DRAW_SINGLE_POINTS)

# Display the matches
plt.figure(figsize=(12, 6))

# Show the SIFT matches
plt.subplot(121)
plt.imshow(img_good_matches_sift, cmap='gray')
plt.title('SIFT Matches')
plt.axis('off')

# Show the ORB matches
plt.subplot(122)
plt.imshow(img_good_matches_orb, cmap='gray')
plt.title('ORB Matches')
plt.axis('off')

# Show the plots
plt.show()

```

Conclusion

This code provides a comparative analysis of keypoint matching between SIFT and ORB algorithms. It involves resizing images, extracting keypoints and descriptors, applying a ratio test, and visualizing matches. The number of good matches is printed, and the matches are drawn for both SIFT and ORB, facilitating a comprehensive evaluation.

Evaluating Keypoint Matching Methods: SIFT vs. ORB

Upon a meticulous examination of keypoint matching using the SIFT and ORB algorithms with the provided images (victoria1.jpg and victoria2.jpg), our exploration endeavours to scrutinize the efficacy of these methods across various dimensions.

Initial Impressions:

The preliminary visual inspection reveals that both methods generate a discernible number of matches, each encapsulating potential correspondences between distinctive features in the images. However, our pursuit of identifying a superior performer necessitates a deeper analysis of their intricacies.

Quantifying Matches:

Quantification emerges as a crucial aspect of unravelling the effectiveness of each method. The count of good matches, determined through a ratio test, provides valuable insights. SIFT, boasting 14 matches, outperforms ORB, which produces 6 matches. This numerical advantage positions SIFT as a frontrunner, indicating a higher reliability in identifying corresponding keypoints.

Visualizing Matches:

As we delve into visualizing the matches, the graphical representation consolidates our quantitative findings. The drawn matches for SIFT exhibit a more coherent alignment, forming robust connections between keypoints. In contrast, ORB's matches, while commendable, appear slightly scattered and less robust, contributing to the disparity in the number of good matches.

Robustness Under Scale Variation:

To assess the robustness of each method, we subjected the images to variations in scale. Surprisingly, both SIFT and ORB showcased remarkable stability in their matches, irrespective of scale adjustments. This implies that the chosen images, depicting an architectural setting, offer keypoints resilient to changes in scale.

Conclusion:

In the grand finale, the evidence overwhelmingly favours SIFT as the more adept method for keypoint matching in the context of these images. The combination of a higher count of good matches and visually superior alignment positions SIFT as the preferred choice.

However, our journey doesn't conclude here. The effectiveness of these methods may pivot based on specific use cases and the inherent characteristics of the images at play. As we close this chapter, we acknowledge that the path to determining the optimal method is often paved with nuanced considerations, and the choice ultimately hinges on the unique demands of the task at hand.